

Frozen supervised decoder vs same raw Jev queries: fresh source-group 95% intervals
No fresh LLM comparison or wall-speed result is available

